Sixth Sense — Attempter Guidelines

Last updated: August 2026

Contents

Select a section below to jump directly to it.

CONFIDENTIAL | apalowodavid9@gmail.com+outlier | 2026-08-27T14:42:09.179Z

Project Objective

Sixth Sense builds a frontier dataset for challenging latent visual reasoning — using media-grounded clues to reach conclusions that are not directly visible or explicitly stated.

Until now this project has collected two kinds of tasks. Complex visual reasoning, where the answer is somewhere in the media and the challenge is finding it — that kind is not for HSS. And latent visual reasoning, where the answer is not in the media at all and a person has to work it out — that kind is good for HSS.

Moving forward we only want the second kind. Every prompt you write should challenge latent visual reasoning, and you select the domain and subdomain that name the reasoning your prompt requires.

For each task you select an image or video from the media library. You then select the domain and subdomain, annotate the evidence window (video only), write the prompt, review the generated model responses, write the rubrics, write the golden answer, and confirm the LLM judge's ratings.

The Two Core Requirements

Every task must satisfy both. They are referred to by name throughout this document rather than restated in full.

Requirement	What it means
R1 — Latent visual reasoning	The primary challenge is reasoning, not perception. A correct answer must compare, infer, predict, rank, track, or explain. Simple perception may serve as evidence inside a larger inference, but never as the whole task.
R2 — Model failure	At least two generated model responses must fail one or more mandatory rubric criteria — and must fail because of the targeted domain–subdomain capability, not because of ambiguity, unclear wording, poor media quality, missing context, or outside knowledge.

One Prompt, One Central Challenge

A prompt may ask one question or several. Multiple questions are allowed only when every question tests the same central visual reasoning challenge from the selected domain and subdomain. Questions may request different outputs and focus on different moments; they may not mix unrelated capabilities, domains or subdomains.

Each question must have one clearly correct answer, and the prompt as a whole must support a single objectively verifiable golden response. Equivalent wording is acceptable; several equally valid answers are not.

The Caption Test

Imagine someone has written a complete caption of your image or video. It lists every object and where it is, what each person looks like and how they are standing, everything that happens in every frame, and every word spoken. It is as long as it needs to be, and it is perfectly accurate.

Would your answer be in that caption?

Yes, it would be in the caption. Not for HSS.

The prompt is asking the model to look something up. Write a different one.

No, it would not be in the caption. Good for HSS.

The description is complete and the answer still is not in it, so a person has to work it out. This is the data we need.

Good for HSS	Not for HSS
Requires a meaningful reasoning step, not a lookup.	Asks only for an obvious count, color, object, text, or clothing attribute.
The conclusion comes from comparing or connecting details in the media.	The answer can be copied from one immediately visible detail.
Uses simple perception as evidence within a broader inference.	Uses recognition or counting as the entire task.
All questions test the same central visual reasoning challenge.	Combines questions that test different domains, subdomains, or capabilities.
Rubrics cover exactly what the prompt asked for.	Rubrics add optional observations the prompt did not request.

The task should be difficult for current frontier models while remaining objectively answerable by a careful person using only the provided media.

Workflow at a Glance

Ten steps, in the order the task form presents them. Video tasks add the evidence window at Step 6.

Step	Action
1. Select the Media	At the prompt step, click the paperclip and attach exactly one image or video from the File Library. Skip anything containing PII, minors, sensitive or graphic content, or NSFW material while you are still browsing.
2. Write the Prompt	Ask one or more focused questions that satisfy R1 and R2. Use the Caption Test before you submit.
3. Select the Domain and Subdomain	Select the domain and subdomain naming the reasoning your prompt requires — the step a caption of this media could not perform.
4. Review the Model Responses	Three responses are generated. Give each a Pass or Fail verdict and a one- to two-sentence justification.
5. Report How Many Models Failed	Enter how many of the three failed. Fewer than two means the prompt is not hard enough — revise it and try again.

Step	Action
6. Set the Media Type and Evidence Window	Declare whether you uploaded an image or a video. Video tasks then require a start and end timestamp for the segment containing the evidence.
7. Write the Rubrics	Write mandatory, equal-weight criteria covering each required answer element.
8. Write the Golden Response	Answer every question and connect the required evidence to each conclusion.
9. Confirm the LLM Judge Rubric Verdicts	For each response, agree or disagree with the judge and justify your judgement in one or two sentences.
10. Confirm the Domain and Subdomain	Confirm the domain and subdomain at the end of the task. This is the selection the task is tagged with, so check it still matches the prompt you finished with.

Step 1: Select the Media

You choose the media for your own task. At the prompt step, click the paperclip below the prompt box and attach exactly one image or video from the File Library before you write anything — the media comes first, and the prompt is written to fit it.

Using the Library

- Click the paperclip below the prompt box, then choose Browse library. The files there are pre-vetted and provided by Outlier.
- Search by keyword. The search runs against the metadata attached to each file, so try objects, settings, and actions — “crane”, “kitchen”, “crossing”.
- Filter by domain if it helps, but treat the domains as suggestions only. A file tagged for one domain will often support a stronger task in another, so look across all of them rather than trusting the tag.
- Switch to thumbnail view using the control on the right of the picker. From there you can see each file properly and play the videos before committing to one.
- Swap the media out whenever it stops working. If the scene will not carry a strong prompt once you start iterating, remove it and select another — do that before you invest time in rubrics.

What Makes a Good Pick

Run the Caption Test on the media itself, before you write anything. If a complete description of the scene would leave nothing to work out, no prompt will save it — go back and choose something else.

Look For	Skip
Depth, occlusion, and things that carry on past the edge of the frame.	Flat, fully visible scenes where everything is already on show.

Look For	Skip
Several agents, or one agent acting on an environment.	A single subject against a plain background.
A visible mechanism, process, or chain of events.	Static scenes with no cause and nothing in motion.
Traces of something that already happened — wakes, marks, spills, disturbance.	Scenes where nothing has happened and nothing is about to.
For video, evidence that develops across the clip.	Clips that are answerable from the opening frame.

When to Choose Video

Images and videos both produce strong tasks, and the library has plenty of each. Video is worth reaching for, though — it gives you more to constrain, more to compare, and more places to hide the answer, and it tends to break models more reliably.

- Social tasks usually land better on video. An intention, a relationship, or a power dynamic shows itself across a sequence of actions and reactions. A single frame gives you an expression and little else, and models read expressions well.
- Temporal and causal tasks benefit for the same reason. Retrodiction, mechanistic causality and extrapolation all get harder when the evidence is spread across a clip instead of sitting in one frame.
- Video lets you add visual constraints. Pin the target down by what they are wearing, where they are standing, and what they have just done. That narrows the question without revealing the answer, and forces the model to track one agent through the clip rather than guessing from a thumbnail.
- Video lets you ask about change over time. What shifted between two moments, how someone behaved after something happened, why a reaction followed what preceded it. These questions have no single-frame equivalent, and they are among the most reliable ways to break a model.
- Give yourself room in the evidence window. Where the clip supports it, work at the top of the 5–10 second range rather than the minimum. A window that spans a genuine change cannot be resolved by sampling a few frames.

Review the Input

Establish what the media actually shows, and what can reasonably be inferred from it, before you write anything.

- Inspect the entire image, including relevant background detail.
- For video, watch the full clip and rewatch important moments.
- Consider event order, movement, expressions, and environmental context.
- Use only what is visible, plus necessary commonsense reasoning.

Use	Unacceptable
Clearly visible details — for video, details spanning different parts of the clip.	Cues that are unreliable or unclear even to a careful human.
Well-supported commonsense inference.	Assumptions the media does not support.
For video, evidence spanning 5–10 seconds or more.	Windows shorter than 1% of the total duration of the video, or questions answerable from the first frame.
Details that can be described unambiguously in words.	References that can be read more than one way.

Content Safety

The platform shows this warning beside the prompt box, and it is on you to honor it. Before you attach anything, ask:

"Does this image or video contain any personally identifiable information (PII), minors where faces are not blurred or in a documentary context, sensitive or graphic content, or NSFW material?"

- Yes — do not build a task on it. Go back to the library and select different media.
- No — attach it and write your prompt.

Step 2: Write the Prompt

Write one or more focused questions that primarily require latent visual reasoning of the kind you intend to challenge (R1), and that are hard enough to break at least two models (R2).

The model should need to compare, infer, predict, rank, track, or explain relationships across details — not merely identify or count what is visible.

A Strong Prompt

- Is natural, realistic, and grammatically correct.
- Names the specific person, object, or event in question.
- Cannot be answered without the image or video.
- Does not refer to audio.
- Is visually focused rather than text focused.
- Contains no timestamps — identify moments by what is visible.
- Does not reveal or assume the answer.
- Names things accurately — every person, object and attribute used to identify your subject is actually there and correctly described.
- Uses an appropriate level of certainty.
- Yields one clearly correct golden response, not several equally valid ones.
- Is challenging but solvable by a careful person from the media alone.

- A careful person can answer within ~5 seconds.
- Any questions have one objective answer that ~99% of careful viewers would give.

Strong Latent-Reasoning Prompt	Weak Simple-Perception Prompt
Will the cup fit inside the box based on both its visible width and height?	How many cups are visible?
Which container can hold more liquid based on the relative width, height, and shape of each?	What color is the container?
Focus on the person walking alone just above the center of the image, nearest the Japanese Rising Sun flag. To reach the “B” sign, should the person turn around or continue in the same direction?	Which direction is the person walking?
After being called to the front, what emotion is the person feeling based on both their facial expression and their immediate body language?	Is the person smiling?

For video tasks, design the question so the answer requires evidence from multiple moments within the evidence window, and so it cannot be resolved by inspecting only the frames immediately around any one moment.

Use the extra room video gives you. Identify the person or object by appearance, position, and what they have just done — stacked visual constraints narrow the target without giving the answer away — then ask about something that changed: what shifted between two moments, how a person behaved after an event, why one reaction followed another. Change-over-time questions have no single-frame shortcut, which is exactly what makes them hard for models.

Step 3: Select the Domain and Subdomain

With the prompt written, select the domain and subdomain naming the reasoning step it requires — the step a person has to perform to get from what is visible to the answer. Your prompt, rubrics and golden response must all stay consistent with the pair you select.

- Choose for the prompt as a whole, not for one sub-question.
- Base the choice on the main inference your prompt requires — not on the objects, events, text, or attributes that happen to be visible in the media. One image or video may support several domains depending on how the task is designed.
- If your prompt does not fit any subdomain, that is evidence it is not for HSS — go back and rewrite it rather than forcing a selection.

The Four Domains and Ten Subdomains

Every prompt that is good for HSS works because of a specific step the caption cannot perform. Those steps are the ten subdomains below, grouped into four domains. If your prompt does not map obviously onto one of them, that is usually a sign it is not for HSS. The example prompts are the current hero samples — approved tasks that broke the models. Treat them as illustrations of the reasoning each subdomain targets, not as prompts to reuse.

Subdomain	What the reasoning requires	Example Prompt
Temporal & Causal Dynamics (time) — Recovering the scene's state at a moment other than the one depicted, or the causal process connecting those moments. The media is one slice; the answer lies behind it, ahead of it, or in the mechanism driving it.		
Retrodiction	Inferring what already happened based on visual traces.	Focus on the man carrying the box. From what side of the frame would he have emerged from if this was a video and he walked into the scene in the same direction and trajectory that he is walking now?
Mechanistic Causality	Identifying the immediate physical mechanism driving an ongoing event. This is about understanding the "how" and "why" of current motion or state changes, not predicting the future.	There was a moment when player number 18, wearing the black jersey, fell right in front of the camera in their own defensive zone. Why did he fall?
Extrapolation	Predicting the immediate kinematic trajectory or long-term systemic outcome of the current scene.	Focus on the orange and black fish that is partially cut off on the far right side of the image. What direction is it traveling if it continues on the same path? Is it going up or down and to the left or right?
Physical & Spatial Logic (space) — Recovering unshown structure of the physical scene: what exists outside view, what properties objects hold beyond their appearance, and how the layout resolves from a coordinate the camera never occupied.		
Hidden & Invisible	Reasoning about entities or physical properties that are not directly visible in a scene. This involves deducing the existence, shape, and location of objects that are occluded, off-screen, or inside containers, as well as inferring non-visual physical attributes (mass, friction, temperature, material rigidity) and invisible forces (wind, pressure, gravity) from observable visual cues.	Consider Balloon 1 as the closest to the camera, Balloon 2 as the second closest, and Balloon 3 as the third closest. Which balloon has the least outward force exerted on its fabric?

Subdomain	What the reasoning requires	Example Prompt
Affordance & Feasibility	Evaluating whether an action, state change, or physical interaction can be successfully achieved given the capabilities of an agent, the properties of objects, and the constraints of the environment. This encompasses both physical compatibility (such as spatial fit) and functional outcomes (such as illuminating a room by toggling a switch).	Will another vehicle be able to park between the red and black sedans?
Spatial "Alien Viewpoint" and Reachability	Mentally re-projecting the 3D scene to calculate line-of-sight or reachability from a different physical coordinate.	From the position of the soldier wearing the yellow helmet closest to the bridge, can he see the yellow shipping container in the background without any person partially blocking his view?
Social Understanding (mind) — Recovering the unobservable mental and normative layer: what agents believe, want, feel, attend to, and fail to know — and what the situation's unwritten rules permit, forbid, or make awkward.		
Theory of Mind	Reasoning about the internal beliefs, knowledge gaps, deceptions, and emotional states of an agent.	Why did the man put the white powder on it?
Social Role, Norm & Power Dynamics	Reasoning about unwritten social contracts, etiquette, and cultural appropriateness in a scene.	Of the 4 doctors in the room, which one is most likely the anesthesiologist? What clue would lead to this conclusion?
Abstract & Contextual Inference (context) — Recovering what a scene implies but never shows: the conventions and circumstances that give it meaning, and the outcomes of worlds it doesn't depict. The answer follows from applying a rule or reading a context, not from observing the frame.		
Change & Consequence	Reasoning about a change the scene doesn't contain: given the change, inferring its consequence; or given the consequence, inferring the change required to produce it.	If I were to hold the blue rotating wheel in the background while the metal ball is sliding down the zigzag path, what would happen to the chain reaction?
Patterns & Pareidolia	Identifying structure in random data or showing specific hardwired bias to detect familiar patterns — especially faces and human figures — in random noise, inanimate objects, or ambiguous shapes.	Based on the labels on the other beakers, what number can be inferred on the label of the last beaker despite being difficult to read?

CONFIDENTIAL | apalowodavid9@gmail.com+outlier | 2026-08-27T14:42:09.179Z

Two More Things an HSS Task Needs

Clearing the Caption Test is necessary but not sufficient. A task that is good for HSS also has to be quick and it has to be shared.

- Not a deliberative search. A person should be able to answer within about five seconds of looking — no enumerating steps, no counting, no arithmetic. If solving it feels like work rather than intuition, it is the wrong kind of hard.
- The intuition has to be shared. Another careful person looking at the same media should land on the same answer. If reasonable people would disagree, the task measures one attempter's opinion rather than a human competence.

The target is a narrow band: easy for a person, impossible to read off the pixels.

Step 4: Review the Model Responses

Review every generated model response and determine which ones fail to answer a question correctly or omit required evidence.

Use those failures to identify each mandatory answer element an ideal response must contain. Each one becomes a separate rubric criterion at Step 7.

Three responses are generated, labelled A, B and C. Each one has to be given a verdict before you can continue.

Pass or Fail. A response fails if it gets any part of the answer wrong, or omits something the prompt explicitly asked for. Judge the final answer only — never the reasoning trace shown above it.

A one- to two-sentence justification. Say what the response got wrong, not just that it was wrong. "The response gets the first question right but the second question wrong" is the level of specificity to aim for.

Step 5: Report How Many Models Failed

The form then asks how many of the three responses failed. Enter the number that matches the verdicts you just gave — the two have to agree.

Fewer than two failures means the prompt is not hard enough. Revise it and regenerate rather than submitting a task that will be rejected. This is R2, and it is the single most common reason work is thrown away.

Step 6: Set the Media Type and Evidence Window

The form asks what type of media you uploaded. Answer Image and you are done with this step. Answer Video and two timestamp fields appear: identify the complete segment containing the evidence needed to solve your prompt, and record its start and end.

Timestamps belong in the Evidence Window only.

> Do not include timestamps in the prompt or in the golden response. Earlier guidance allowed this — for example, “the five friends pose in front of the bus at approximately 00:35…” in a prompt, or “from 00:36–00:41 the man scratches his neck” in a golden answer. That is no longer acceptable.
>
> Identify moments by what is visible instead: “when the five friends pose in front of the bus” rather than “at 00:35.” The Evidence Window field is the single place timestamps appear.

Rules

- Record both a start and an end value in XX:XX format (minutes:seconds).
- Target a window of 5–10 seconds that captures the full action, sequence, comparison, or change being tested. Shorter windows are allowed when the evidence genuinely occupies less time, but the minimum is no less than 1% of the video’s total duration — anything under that is not acceptable, and never annotate a single frame or instant.
- Prefer the longer end of that range. A window that spans a real change — a shift in behaviour, a sequence of actions and reactions, a mechanism running its course — is harder than a short one, because it cannot be resolved by sampling a few frames. Social and temporal tasks in particular tend to need the room.
- Design the prompt around evidence appearing later in the clip whenever the video supports it. Anything answerable from the first few frames does not meaningfully test video understanding and is not acceptable.
- When evidence appears in separate parts of the clip, record the earliest relevant timestamp as the start and the latest as the end, even though the moments are not continuous.

Situation	Record
One continuous segment	00:24–00:33
Evidence at 00:18 and again at 00:47	00:18–00:47

Step 7: Write the Rubrics

Rubrics define the final-answer requirements a response must satisfy to be fully correct. Because a prompt may contain multiple questions, the rubric set must capture every mandatory answer element for every question.

Every item is mandatory and equally weighted — no points, weights, or partial credit. Rubrics score final-response content only, no reasoning or justifications unless the prompt explicitly requests it. A golden response must satisfy every item.

Design Principles

Principle	Requirement	Meaning
Equal treatment	No weights or point values.	All items are evaluated equally; failing any one means the response fails.
Mandatory	Include every — and only — required final-answer element.	No optional observations, bonus details, or unrequested explanations.
Objective	Each item is verifiable as met or not met.	Do not judge style, confidence, or hidden reasoning.
Atomic	One requirement per item.	Split any criterion joined by "and."
Self-contained	Name the relevant person, object, event, or conclusion.	Avoid "it," "they," "this," or "the others."

Rules

- Start each item with a verb, with "response" as the implied subject: "States…", "Identifies…", "Mentions…", "Provides…"
- Include the final-answer elements the prompt requires — every one of them, and nothing else. No criteria for optional observations, helpful context, or reasoning/justifications unless the prompt explicitly requests it. Meaning if the prompt says "and explain why?" then criteria can be included for reasoning.
- Write one requirement per rubric; split anything that combines requirements with "and."
- Make each rubric self-contained by naming the relevant person, object, or event.
- Avoid duplicate or overlapping criteria.
- Do not write a criterion that cannot be verified from the visuals alone — for example, one that depends on audio.
- Evaluate meaning rather than exact wording, and correctness rather than presentation style or confidence.

Example: if a prompt asks whether a cup fits in a box based on width and height, write separate criteria for the fit conclusion, the width comparison, and the height comparison. Do not add criteria about color, material, or other unrequested details.

Then Verify R2

Compare each model response against every rubric criterion and confirm that at least two responses fail one or more of them, and that those failures are caused by the targeted capability in the selected domain and subdomain. If fewer than two fail, revise the prompt — and the rubrics if needed — then regenerate and review the responses before writing the golden answer.

Step 8: Write the Golden Response

The golden answer is the verified reference response used to score model answers. It must answer every question directly and satisfy every rubric criterion without unnecessary repetition.

Note: It is acceptable to copy a passing model response to use as your golden response, but it must be edited significantly enough to sound natural. Tasks that contain golden responses that have not been edited will not pass.

Recommended Structure

Part	What to Include
Conclusion	State the direct answer immediately.
Evidence	Identify every required visual, temporal, textual, or commonsense cue.
Reasoning	Explain how those cues support the conclusion.
Qualification	Note uncertainty only when the evidence does not support a fully certain answer.

Grounding and Certainty

Type	Example
Direct observation	The person repeatedly looks toward the exit and steps backward.
Supported inference	These actions suggest the person may be preparing to leave.
Unsupported statement	The person has decided to leave because they are angry.

Example

Prompt: Based on the driver's position and the parked vehicle's obstruction, can the driver currently see the pedestrian?

Golden answer: No. The parked vehicle sits directly between the driver and the pedestrian, blocking the driver's line of sight. The pedestrian is positioned within the area occluded by the vehicle and would become visible only after moving farther beyond it.

Step 9: Confirm the LLM Judge Rubric Verdicts

After the golden response is submitted, the LLM judge evaluates each model response against your rubric criteria and returns a verdict for each one.

For each response, decide whether you agree or disagree with the judge's verdict, and give a one- to two-sentence justification for your judgement. Refer back to the Pass/Fail verdicts you gave at Step 4 — the two should tell the same story. Judge only the response's final-answer content, never its reasoning trace.

A written justification is required here, unlike the verdicts at Step 4 where it is a short note. Evaluate meaning rather than exact wording: a criterion reading "States that the ladder's feet rest on gravel rather than concrete" is satisfied by "the base is standing on loose stones".

CONFIDENTIAL | apalowodavid9@gmail.com-outlier | 2026-08-27T14:42:09.179Z

Step 10: Confirm the Domain and Subdomain

The final step asks you to confirm the domain and subdomain. Whatever is recorded here is what the task is tagged as, so treat it as the selection that counts rather than a formality.

Re-read your prompt and check the pair still describes the reasoning it actually requires. Prompts move while you work on them — a question that started as a sightline problem often ends up turning on a change to the scene, or the other way round — and the tag needs to follow the prompt you finished with, not the one you set out to write.

Confirm your rubrics and golden response are consistent with it too. If the pair no longer fits, change it here.

⚠️ Common Errors to Avoid

The errors causing the most rejections right now. Each is a failure of a requirement already covered above. The sections below show the real rejected tasks behind each one — click Image or Video to open the media.

Common Error	Why the Task Gets Rejected	How to Avoid It
Ambiguous Prompts	More than one answer is defensible, so there is no single golden response. The difficulty is in the wording, not the reasoning.	Name one target and one thing to determine about it. Cut vague comparatives (“substantially repositioned,” “fully hidden”). Would ~99% of careful viewers give your exact answer?
Perception-based Prompts	The answer is a lookup an exhaustive caption would already contain. Fails R1 however hard the models find it.	Run the Caption Test first. The prompt must force a compare, infer, predict or retrodict step; perception can only be the evidence inside it.
Non-Mandatory Rubric Criteria	Criteria for reasoning or detail the prompt never asked for, or near-duplicates. Model failures on those items do not count towards R2.	One criterion per element the prompt explicitly asks for, and nothing else. No “explain why” in the prompt means no explanation criterion.
Video Solvable From a Single Frame	A still from anywhere in the evidence window answers it. The clip does no work and the task fails as a video task.	Build the question around something that changes. The answer should need evidence from at least two separate moments.
Subdomain Does Not Match the Reasoning	The tag was chosen from what is visible rather than the inference required. Patterns & Pareidolia and Hidden & Invisible are the most misapplied.	Tag the step a person performs to reach the answer, then re-check the pair at Step 10.

CONFIDENTIAL | apalowodavid9@gmail.com-outlier | 2026-08-27T14:42:09.179Z

Common Error	Why the Task Gets Rejected	How to Avoid It
Emotion Treated as Automatic Ambiguity	Valid Theory of Mind tasks reworked or rejected on ambiguity grounds. Some interpretation is inherent to reading emotion.	If the reading is correct and the vast majority of careful viewers would agree, the task stands.

1. Ambiguous Prompts

Media	Rejected Task	Why it Fails
[Video]	**Prompt:** If the smoke cloud between the tanks continued expanding in the same area, which visible tank would become fully hidden from the camera first?	Hypothetical and unbounded. "Continued expanding" sets no rate or direction and "fully hidden" has no threshold, so different viewers pick different tanks and all can justify it.
[Image]	**Prompt:** Based on the snake's current position and the way it is held, is the nearby person within the snake's immediate reachable path, or would the snake need to be substantially repositioned before it could make contact with its mouth? What visible spatial cues support your conclusion?	"Immediate reachable path" and "substantially repositioned" are undefined. The answer turns on where that line sits, and the image does not fix it.
[Image]	**Prompt:** For the two bags of apples on the top shelf, how is the bag that has the most apples visible through the packaging standing up like that? Why isn't the other one positioned the same way?	The target is identified by a count made through packaging, so viewers disagree on which bag is meant. The "how" then admits several explanations the image never settles.

Important: Do Not Misread Emotion as Ambiguity.

Ambiguity and subjectivity should not be misconstrued in Theory of Mind and emotion-related tasks. There is always some degree of ambiguity with emotion. If the task appears correct and the vast majority of people would agree with it, do not reject it — this is an issue we have seen.

2. Perception-Based Prompts

Media	Rejected Task
[Image]	**Prompt:** What can be seen from the camera screen on the left of the photo?
[Video]	**Prompt:** During the lecture, with "Introductions" displayed on the presentation screen, if the person seated in the last row is asked to speak to the person sitting closest to him, identify the closest person's position relative to him by specifying both the row and the number of chairs away.

Media	Rejected Task
[Image]	**Prompt:** Based on the positions of the people at the table, who is the person holding people's attention by engaging with paper documents, and why?

3. Non-Mandatory Rubric Criteria

Media	Rejected Task	Why it Fails
[Image]	**Prompt:** Assume the background hills are North, foreground is South, left edge of the photo is West and right edge is East. If the runner stands on the start numbers 1-4 facing forward down the track. Is the central green grass field to their left or right? Which cardinal direction will they turn when following the curve ahead? **Rubric as written:** 1. States that the central green grass field is to the runner's left. 2. States that the runner will turn North (or to the left) when following the curve ahead down the track. 3. Explains that standing on the numbers aligns the runner facing East, placing the infield to their left and the upcoming turn toward the North.	Criteria 1 and 2 are the two answers asked for. Criterion 3 grades the reasoning behind them, which the prompt never requested.
[Video]	**Prompt:** Does the lady with the long dark hair have her right hand free? Explain why based on her posture and hand position. **Rubric as written:** 1. States that the lady does not have her right hand free. 2. States that her right arm is extended behind her as she walks along as if she is pulling something. 3. Describes that her posture is slightly angled forward as she walks, as if she is pulling something.	Criteria 2 and 3 are the same observation written twice — both end on "as if she is pulling something." Two requirements dressed up as three.
[Video]	**Prompt:** In the clip one man standing higher up is poked by one man lower down, what is the reason for this? **Rubric as written:** 1. States that the smiling interaction between the two men before the poke signifies friendship.	The prompt asks one thing. All three say the poke was friendly, with 1 and 2 grading the supporting evidence. Criterion 3 alone covers it.

CONFIDENTIAL | apalowodavid9@gmail.com-outlier | 2026-08-27T14:42:09.179Z

Media	Rejected Task	Why it Fails
	2. Identifies that as both people laugh after the poke occurs, the poke was friendly. 3. Identifies the poke as a jokey gesture.	

4. Task Could Be Solved from a Single Frame

Media	Rejected Task	Why it Fails
Video	**Prompt:** Based on the cars parked on the opposite side of the road, as seen from the camera, is there enough space between the silver car and the pickup truck to park a bicycle perpendicular to the curb? Is there enough space between pick-up truck and the white car to park two bicycles side by side, both perpendicular to the curb and not in a single line? **Golden response:** No, a bicycle cannot fit parked perpendicular between the silver hatchback and the pickup truck because the gap is too narrow even for a bicycle's handlebar width. Yes, two bicycles can fit side-by-side and perpendicular between the pickup truck and the white station wagon because that gap is visibly wider and leaves room for two handlebar widths.	The cars are static. Both gap comparisons resolve from one frame, so the evidence window adds nothing.
Video	**Prompt:** In the scene where a woman wearing a yellow-and-green top is seated at a desk with three people sitting in chairs in front of her, if the black speaker mounted high on the left wall were to detach and fall straight down, whose head would be directly beneath its path? Justify your answer using the speaker's position relative to the people in the room. **Golden response:** Nobody. The black speaker is mounted on the far-left side of the room, while the woman in the yellow-and-green top and the three seated people are positioned farther to the right. If the speaker fell straight down, its vertical path would not align with anyone's head.	Speaker fixed, people seated. The drop line reads off any single frame, and the golden response cites no evidence from a second moment.
Video	**Prompt:** Will another vehicle be able to park between the red and black sedans? Explain why. **Golden response:** Yes, another vehicle will be able to park between the red and black sedan. As other	A static gap judgement. The golden response is also weak — “there should be an extra space” infers from the other cars rather than the gap, so

Media	Rejected Task	Why it Fails
	vehicles in the parking line are close to each other, there should be an extra space between the red and black sedans.	it is ambiguous as well as unnecessarily a video.

5. Incorrect Subdomain Selected (Specific Issues with Patterns & Pareidolia and Hidden & Invisible)

Media	Rejected Task	Why it Fails
Image	**Tagged:** Patterns & Pareidolia **Prompt:** Judging from the current positions and orientations of the players on the field, which of the teams between the light colored or the dark colored team is currently in possession? Also, which of the teams is defending the left goalpost?	Should be Social Role, Norm & Power Dynamics. Reading possession and which side a team defends is inference about roles from body orientation, not recognition of a visual pattern.
Video	**Tagged:** Hidden & Invisible **Prompt:** If the conductor turned to face the yellow flag, which instrument groups would be unable to clearly see his conducting gestures?	Should be Spatial "Alien Viewpoint" and Reachability — it projects sightlines from a hypothetical orientation, nothing is concealed. Also subjective: "clearly see" has no cut-off.

Final Checklist

Good for HSS — Check This First

- ☐ The answer would NOT appear in an exhaustive caption of the media.
- ☐ No sub-question would fail the Caption Test on its own.
- ☐ The prompt maps onto one of the ten subdomains, and that subdomain is selected.
- ☐ A person can answer in about five seconds, without counting or arithmetic.
- ☐ Another careful person would give the same answer.
- ☐ The task is not built on counting, ordering, reading, measuring, gaze, or on-screen motion.

Core Requirements

- ☐ R1 — The challenge is an HSS inference: simulating, projecting, attributing or retrodicting. Any simple-perception step serves as evidence within it.
- ☐ R2 — At least two model responses fail one or more mandatory rubric criteria, because of the targeted capability rather than ambiguity, poor media quality, trick wording, or outside knowledge.

☐ If fewer than two models failed, the prompt and/or rubrics were revised and the responses regenerated.

☐ Gave every response a Pass or Fail verdict with a short justification, and reported a fail count that matches those verdicts.

Safety and Selection

☐ Selected media from the library that supports an inference a full description would not contain.

☐ Answered the PII, minors, sensitive/graphic, and NSFW screening question.

☐ Selected different media from the library if any listed issue was present.

☐ Selected the domain and subdomain matching the primary capability tested, not just the objects visible in the media.

☐ Prompt, rubrics, and golden answer all align with the selected domain and subdomain.

☐ Domain and subdomain confirmed at the end of the task, and still matching the prompt as written.

Evidence Window (Video Tasks)

☐ Both a start and an end timestamp are recorded, in XX:XX–XX:XX format.

☐ The window is at least as long as 1% of the video's total duration and targets 5–10 seconds wherever the video supports it.

☐ The task requires evidence from beyond the first few frames.

☐ The earliest and latest relevant timestamps are used when evidence is spread across the clip.

Prompt

☐ Asks one or more focused, natural, grammatically correct questions.

☐ A careful person can answer within ~5 seconds.

☐ Every question tests the same central HSS challenge.

☐ Cannot be answered from the prompt text alone, and does not rely on audio.

☐ Contains no timestamps.

☐ Does not reveal or assume the answer, and uses appropriate certainty.

☐ Does not use counting, ordering, reading, colour, object naming, size comparison, gaze, or on-screen motion as the primary challenge.

☐ Produces one objectively verifiable golden response ~99% of careful viewers would give.

Rubrics

☐ Reviewed every model response before writing rubrics.

☐ Each item starts with a verb with "response" as the implied subject.

☐ Covers every mandatory answer element — and nothing optional.

☐ Addresses answers, not reasoning process or justification unless explicitly requested by the prompt.

☐ Each item is atomic, objective, and self-contained; nothing is joined by "and."

- ☐ No negative-exclusion criteria, no overlapping criteria, no weights or point values, and no grading of presentation style.
- ☐ No criterion depends on audio or anything not verifiable from the media.

Golden Answer

- ☐ States the conclusion immediately and answers every question.
- ☐ Includes all required evidence and covers every rubric criterion without unnecessary repetition.
- ☐ Uses only grounded visual and temporal evidence, at a certainty the media supports.
- ☐ If you utilized a passing model response and pasted it as your golden answer, it must be edited significantly enough to sound natural. Golden responses that are not edited at all will not pass.

LLM Judge Ratings

- ☐ Checked each judge verdict against the corresponding rubric criterion and model response, judging meaning rather than wording.
- ☐ Agreed or disagreed with each judge verdict and justified the judgement in one or two sentences.

Overall

- ☐ A careful person can solve the task from the provided media alone.
- ☐ Difficulty comes from an HSS inference, not from ambiguity, missing context, poor input quality, or simple perception.

CONFIDENTIAL | apalowodavid9@gmail.com+outlier | 2026-08-27T14:42:09.179Z

Appendix A: Question Templates

The prompts below are shapes, not stock. They exist to show the kind of question each subdomain tends to produce. Do not drop one into a task as written — a template used verbatim will almost always be too vague to have one correct answer, or answerable straight off the pixels.

Write every prompt for the specific media in front of you. Name the actual person, object, or moment, and build the question around what that particular scene makes inferable. Read the list for inspiration, then close it. If nothing here fits your media, that is the list's limitation, not your media's — write your own question and check it against the Caption Test.

Subdomain	Question shapes to adapt — never to reuse as written
Retrodiction	What happened to the object, person or animal previously that led to what is shown? Where was the object, person or animal previously? When did something happen — something not visible in the media, but recoverable from visual evidence in it? Who did the thing that led to what is shown? How did the person or animal do what is shown?
Mechanistic Causality	Which object, person or animal is causing what is shown? Why is what is shown happening? The cause may be an object, a person, a force, temperature, pressure, and so on. What allows the object, person or animal to do what it is doing? What differs about the same object at two different moments, and what accounts for the difference?
Extrapolation	What will happen to the object, person or animal if the scene continues? Where will the object, person or animal end up? When will something happen — something not visible in the media, but predictable from visual evidence in it? Who is most likely to do a particular thing next? How would the object, person or animal go about doing a particular thing?
Hidden & Invisible	What is inside or behind the thing that conceals it — a bus, a train, a building? Which object, person or animal has the most or the least of a property: force, pressure, temperature, mass? Where is the camera most likely positioned? Across two similar moments, what is hidden or missing?
Affordance & Feasibility	Can the object, person or animal do a particular thing? Can two agents work together to do it? How could the object, person or animal do it using something visible in the scene? When could the object, person or animal do it?

CONFIDENTIAL | apalowodavid9@gmail.com+outlier | 2026-08-27T14:42:09.179Z

Subdomain	Question shapes to adapt — never to reuse as written
Spatial "Alien Viewpoint" and Reachability	Can the person or animal reach a particular thing? From a particular agent's perspective, how do the objects relate to one another? Add conditions to sharpen it — left to right, furthest to closest, biggest to smallest.
Theory of Mind	Why did the person or animal do a particular thing? Why is the person or animal feeling what they are feeling — happy, sad, jealous, angry, cautious? What does the person or animal not know that led to what is shown? Does the person or animal know a particular thing? What does the person or animal think about a particular thing? What is the purpose of someone doing a particular thing? Someone whispers or speaks — what could they be saying, where the media gives enough to infer it?
Social Role, Norm & Power Dynamics	What is the relationship between the people or animals in the media? Which person or animal is senior to, or leading, another? Who is getting the most attention? Why is something funny — what is the punchline? Why is the person doing a particular thing, where there is sarcasm or irony in it? What should a person do next as part of a cultural tradition? Is what is happening socially acceptable? What is the social intention behind someone's behaviour? Across the whole video, what is the critical turning point if there is a twist?
Change & Consequence	In order to bring about a particular outcome, what would need to change?
Patterns & Pareidolia	What is the pattern — how are the people lined up, what is the order of the colors used? What comes next in the pattern — the next number, the next color? The person or animal did a particular thing; why did they do it, where the reason is a misreading of a symbol, shape or pattern? What do the symbols, icons or marks in the scene mean?

Appendix B: Good Task Examples

One approved example per subdomain — ten in total. Each is a current hero sample; the Model Failures row records how many of the three models the prompt broke. These are illustrations of what each subdomain rewards, not prompts to adapt.

Good Example 1: Retrodiction

Domain & Subdomain	Temporal & Causal Dynamics (time); Retrodiction
Media	https://static.remotasks.com/uploads/6785a5a0d930585c43561838/b12_554.png
Model Failures	Failed all three models
Prompt	Focus on the man carrying the box. From what side of the frame would he have emerged from if this was a video and he walked into the scene in the same direction and trajectory that he is walking now?
Golden Response	Looking at the body orientation and stride of the man carrying the box, he's moving toward the background of the frame, away from the camera and slightly towards his own right. If this was a video and he walked into the scene in the same direction and trajectory that he is currently walking, he would have entered from the left side of the frame, walking across from left to right as he walked past the camera.
Why It Works	His entry into the scene is never shown. The answer comes from reading his body orientation and stride, then running that trajectory backwards past the edge of the frame — a step no description of his posture can perform.

Good Example 2: Mechanistic Causality

Domain & Subdomain	Temporal & Causal Dynamics (time); Mechanistic Causality
Media	https://static.remotasks.com/uploads/6837b01447c25ec5228133dc/b6_97.mp4
Model Failures	Failed all three models
Prompt	There was a moment when player number 18, wearing the black jersey, fell right in front of the camera in their own defensive zone. Why did he fall?
Golden Response	Player #18, wearing the black jersey, fell because he ran into his teammate, Player #22, from behind and lost his balance.Player #18 (black jersey) was skating hard back into his own end, near the goal, and as he tried to make a hard turn into the field to cover the play in front of his net, he ran into his teammate, Player #22, and he lost his balance. He went down flat on the ice right in front of the camera, though Player #22 did not fall. Player #18 got right back up and rejoined the play, with no penalty called and no stoppage, confirming that there was no foul.
Why It Works	The fall is plainly visible; the cause is not. Answering it means tracking where the other players were and ruling out contact the caption would list but not explain.

CONFIDENTIAL | apalowodavid9@gmail.com+outlier | 2026-08-27T14:42:09.179Z

Good Example 3: Extrapolation

Domain & Subdomain	Temporal & Causal Dynamics (time); Extrapolation
Media	https://static.remotasks.com/uploads/6837b01447c25ec5228133dc/b1_2074.jpg
Model Failures	Failed all three models
Prompt	Focus on the orange and black fish that is partially cut off on the far right side of the image. What direction is it traveling if it continues on the same path? Is it going up or down and to the left or right?
Golden Response	The partially visible orange-and-black fish is traveling up and to the right. This is visible from the angle of its body sloping upward and to the right in the image.
Why It Works	A still image with no motion to observe. The heading has to be read off the angle of the body and projected forward past the edge of the frame.

Good Example 4: Hidden & Invisible

Domain & Subdomain	Physical & Spatial Logic (space); Hidden & Invisible
Media	https://static.remotasks.com/uploads/6785a5a0d930585c43561838/b1_4993.png
Model Failures	Failed two of three models
Prompt	Consider Balloon 1 as the closest to the camera, Balloon 2 as the second closest, and Balloon 3 as the third closest. Which balloon has the least outward force exerted on its fabric?
Golden Response	Balloon 1 appears to have the least internal pressure because its fabric has the most pronounced wrinkles and loose folds, indicating less outward tension than the smoother fabrics of the other balloons.
Why It Works	Internal pressure is not a visible property. It is inferred from how wrinkled or taut each envelope is, and then compared across three balloons.

Good Example 5: Affordance & Feasibility

Domain & Subdomain	Physical & Spatial Logic (space); Affordance & Feasibility
Media	https://static.remotasks.com/uploads/6837b01447c25ec5228133dc/b6_302.jpeg
Model Failures	Failed all three models
Prompt	Will another vehicle be able to park between the red and black sedans?

Domain & Subdomain	Physical & Spatial Logic (space); Affordance & Feasibility
Golden Response	Yes, another vehicle will be able to park between the red and black sedan. As other vehicles in the parking line are close to each other, there should be an extra space between the red and black sedans.
Why It Works	The gap is visible; the verdict is not. Answering it means holding a vehicle length in mind and measuring it against the spacing the rest of the parked line establishes.

Good Example 6: Spatial "Alien Viewpoint" and Reachability

Domain & Subdomain	Physical & Spatial Logic (space); Spatial "Alien Viewpoint" and Reachability
Media	https://static.remotasks.com/uploads/6785a5a0d930585c43561838/b12_551.png
Model Failures	Failed all three models
Prompt	From the position of the soldier wearing the yellow helmet closest to the bridge, can he see the yellow shipping container in the background without any person partially blocking his view?
Golden Response	Yes, he can see the yellow container without anyone partially blocking his view. This is because from his position, the soldier has a clear line of sight to the yellow shipping container in the background.
Why It Works	A caption lists every soldier and every container. It cannot resolve the sightline, which has to be rebuilt from a coordinate the camera never occupied.

Good Example 7: Theory of Mind

Domain & Subdomain	Social Understanding (mind); Theory of Mind
Media	https://static.remotasks.com/uploads/6785a5a0d930585c43561838/b1_3374%20(1).mp4
Model Failures	Failed all two models
Prompt	Towards the end of the video, the older gentleman in the dark suit opens a door; what is he most likely feeling after opening it, and why?

CONFIDENTIAL | apalowodavid9@gmail.com+outlier | 2026-08-27T14:42:09.179Z

Domain & Subdomain	Social Understanding (mind); Theory of Mind
Golden Response	**Most likely emotions:** Towards the end of the video (around 23:10), when the older gentleman in the dark suit opens the laboratory door, he is most likely feeling puzzled and shocked. **Why he felt this way:** Just moments earlier (around 22:15–22:25), he personally escorted his visitor (the man in the grey suit) out into the hallway and asked him to sit on the bench to wait. Later when he opens the door expecting to find his guest waiting, the bench is completely empty. The visitor has vanished into thin air without making a sound or leaving a trace.
Why It Works	The caption records that he opens a door and describes his face. It cannot record what he expected to find. Answering requires linking two moments that are minutes apart — he escorts his visitor into the hallway and tells him to wait on the bench, and later opens the door onto an empty bench. The emotion follows from the gap between what he believed and what is there. Nothing in the closing frames explains it, so a model reading only the moment of the reaction has no way to reach the answer.

Good Example 8: Social Role, Norm & Power Dynamics

Domain & Subdomain	Social Understanding (mind); Social Role, Norm & Power Dynamics
Media	https://static.remotasks.com/uploads/6785a5a0d930585c43561838/b1_3724%20(1).mp4
Model Failures	Failed all three models
Prompt	In this video, two women (one with short hair and the other with long hair) are sitting on a sofa and talking. Which of them appears to be more senior (disregarding age)? Can you tell this from the photo of them taken in the video?
Golden Response	Yes, the woman with short hair holds a more senior position than the woman with long hair. This is evident from their positioning in the group photo: the woman with short hair stands at the center of the group (specifically, in the exact middle of the second row), whereas the woman with long hair is kneeling at the center of the first row. According to the conventions of official group photos, individuals with greater seniority typically stand closer to the center of the formation. Therefore, the woman with short hair holds higher seniority.

Domain & Subdomain	Social Understanding (mind); Social Role, Norm & Power Dynamics
Why It Works	A caption can describe both women, where each sits, and exactly where each stands in the group photo. It cannot say which of them outranks the other. The answer comes from applying a convention to an arrangement — in a formal group photo the senior figure occupies the centre of the formation — and the evidence sits in a different part of the clip from the conversation the prompt opens with, so the two have to be connected. Models read the faces and the sofa and never think to consult the photo. The prompt also closes off the obvious shortcut by ruling out age, so seniority cannot be guessed from who looks older.

Good Example 9: Change & Consequence

Domain & Subdomain	Abstract & Contextual Inference (context); Change & Consequence
Media	https://static.remotasks.com/uploads/6837b01447c25ec5228133dc/b5_457.mp4
Model Failures	Failed all three models
Prompt	If I were to hold the blue rotating wheel in the background while the metal ball is sliding down the zigzag path, what would happen to the chain reaction?
Golden Response	Once the metal ball is on the zigzag path, holding the blue spinning wheel would not affect the chain reaction because it is evident that the ball has already passed that part of the mechanism.
Why It Works	The intervention never happens. Answering it means simulating a change the media does not contain, then recognising that by the moment named the wheel is already behind the ball — so the consequence is no consequence at all.

Good Example 10: Patterns & Pareidolia

Domain & Subdomain	Abstract & Contextual Inference (context); Patterns & Pareidolia
Media	https://static.remotasks.com/uploads/6837b01447c25ec5228133dc/b6_573.mp4
Model Failures	Failed all three models
Prompt	Based on the labels on the other beakers, what number can be inferred on the label of the last beaker despite being difficult to read?

CONFIDENTIAL | apalowodavid9@gmail.com+outlier | 2026-08-27T14:42:09.179Z

Domain & Subdomain	Abstract & Contextual Inference (context); Patterns & Pareidolia
Golden Response	The final beaker in the row can be inferred to read "80" as the beginning of the clip shows there are clearly 5 beakers, and the first starts at "40," while the next 3 rise in increments of 10 to read "50," "60," and "70." Therefore, it can be assumed that the fifth beaker would follow the same pattern and increase by 10.
Why It Works	The final label is never legible, so the caption can only record it as an unreadable mark. The answer comes from recognising the series the other labels establish and carrying it forward one step.

Appendix C: Bad Task Examples

Bad Example 1: Simple Perception, No Inference

Domain & Subdomain	Physical & Spatial Logic; Hidden & Invisible
Media	https://static.remotasks.com/uploads/6785a5a0d930585c43561838/b12_327.png
Prompt	The photographer of that picture has the attention of only one person in that scene. What is the colour of their top?
Reviewer verdict	simple perception - "what color is it" (CB disabled)
Why it fails	The answer is a colour. Even the setup — that one person is looking at the photographer — is something a caption states outright. Nothing has to be worked out. Attaching it to Hidden & Invisible does not change what the question asks for.

Bad Example 2: No Determinate Answer

Domain & Subdomain	Physical & Spatial Logic; Hidden & Invisible
Media	https://static.remotasks.com/uploads/6785a5a0d930585c43561838/b1_839.png
Prompt	None of the structure holding these baskets up is visible anywhere in this picture. What shape must that structure have?
Reviewer verdict	Thank you for the attempt. This task sadly fails because the prompt is ambiguous and there is no clear GTFA, as the responses clearly demonstrated. There is no obvious shape that the structure must have; it could be many different shapes, and the criteria do not actually mention any specific shape. The structure underneath the baskets could be many, and there are several different ways one could raise the

Domain & Subdomain	Physical & Spatial Logic; Hidden & Invisible
	baskets. Please ensure that your future prompts are more specific and unambiguous.
Why it fails	The instinct is right — the support is genuinely invisible and has to be inferred. But there is no single shape it must have; several are consistent with the image, and the criteria never commit to one. If the models can all give different defensible answers, there is no ground truth to fail against.

Bad Example 3: Presupposes Answer

Domain & Subdomain	Social Understanding; Social Role, Norm & Power Dynamics
Media	https://static.remotasks.com/uploads/6837b01447c25ec5228133dc/b6_421.mp4
Prompt	Why in the course of this game does the first player to be hit by the blue-shirted girl (carrying the red ball) seem to have such an immediate and consequential response as compared to the second player she hits with that ball who, conversely, rather gives an almost apologetic and polite *non*-reaction?
Reviewer verdict	The prompt itself is rather unnatural sounding and has a very technical sound to it. More importantly, it presupposes the conclusion about what the model should infer from the 2 players reaction. A simpler prompt would be "The girl in the blue shirt hits two different players with the red ball during the game. Why does the first player she hits react so differently from the second?" Looking at the quadball rulebook, there is no penalty during friendly fire, so there is an issue whether the models actually failed. This also makes your criteria 3 factually incorrect. Criteria 5 is not necessary to answer the prompt request.
Why it fails	Describing one reaction as “immediate and consequential” and the other as “almost apologetic” tells the solver what to conclude before they have looked. The phrasing is also unnatural. “The girl in the blue shirt hits two players with the ball. Why do they react so differently?” asks the same thing and gives nothing away.

Bad Example 4: Details Not in Media

Domain & Subdomain	Temporal & Causal Dynamics; Mechanistic Causality
Media	https://static.remotasks.com/uploads/6785a5a0d930585c43561838/b1_906.png
Prompt	Imagine a person standing centered immediately in front of the silver pedestal, with their back to the altar and facing the camera. Without changing position, they turn ninety degrees to their left. Which red floor speaker are they now facing more

CONFIDENTIAL | apalowodavid9@gmail.com-outlier | 2026-08-27T14:42:09.179Z

Domain & Subdomain	Temporal & Causal Dynamics; Mechanistic Causality
	directly: the one on the viewer's left or the viewer's right? They then turn another one hundred and eighty degrees in place. Which of the two lower niche statues are they now facing more directly: the one on the viewer's left or the viewer's right? Answer both using viewer-left/viewer-right, not the person's left/right.
Reviewer verdict	This is a bad task. There is no "silver pedestal" in the provided image. All color hues are gold. And the existence of "red floor speakers" are also questionable. None of the model responses has failed.
Why it fails	There is no silver pedestal in the image — everything is gold — and the red floor speakers the prompt turns on are questionable too. A prompt built on objects that are not there cannot have a verifiable answer, and unsurprisingly no model failed. Check every object you name against the frame before you submit.

Bad Example 5: Timestamps in Prompt

Domain & Subdomain	Temporal & Causal Dynamics; Mechanistic Causality
Media	https://static.remotasks.com/uploads/6837b01447c25ec5228133dc/b13_597.mp4
Prompt	Evidence Window:00:05 - 00:14The ball rolls almost the full length of the court after a single putt. What property of the surface it travels across is primarily responsible for it covering that distance ?
Reviewer verdict	perception-based prompt + including timestamps in the prompt.CB disabled (this is their first project on outlier)
Why it fails	The prompt opens by pasting in the evidence window. Timestamps belong in that field and can be in the golden response, nowhere else — identify the moment by what is visible instead. The underlying question is also perception: the surface property is stated by the scene rather than inferred from it.

Bad Example 6: Contrived

Domain & Subdomain	Temporal & Causal Dynamics; Extrapolation
Media	https://static.remotasks.com/uploads/6785a5a0d930585c43561838/b1_695.png
Prompt	Consider the two large fish in the image with black coloring or marks. Call the one on the left Fish A and the one on the right Fish B. If the fish tank was infinitely large (so the camera was in the water) and the camera extended left and right infinitely,

CONFIDENTIAL | apalowodavid9@gmail.com+outlier | 2026-08-27T14:42:09.179Z

Domain & Subdomain	Temporal & Causal Dynamics; Extrapolation
	which fish would reach the position of the camera first if they swam straight in their current trajectory?
Reviewer verdict	Contrived
Why it fails	An infinitely large tank with the camera submerged and the frame extending forever is not a scenario the media supports, and nobody would ask it. Conditions should be changes a person could picture happening in that scene, not geometry problems bolted onto it.

CONFIDENTIAL | apalowodavid9@gmail.com+outlier | 2026-08-27T14:42:09.179Z